

Two for Two: Team Essay 2 Group 2

Jonathan Ferdinand, Devina Gera, Archimedes Li, Henry Liu,
Katherine Shi, Sam Stevens

March 2025

All code can be found on our Github repository at https://github.com/mshki/math-456/tree/main/essay_2, or in Appendix A

1 Introduction

Our model is a multiple linear regression model that predicts the number of sales a product would get based on budgets for various methods of advertising. Multiple linear regression is a function of multiple features, X_1, X_2 , that predicts a label \hat{y} in the form of $\hat{y} = \beta_1 X_1 + \beta_2 X_2 + \beta_0$. The benefits of using a linear regression are its simplicity and high level of explainability. This makes it a good candidate for modeling a simple relationship between variables. However, a drawback to linear regressions is their low complexity. We selected a dataset that fits these constraints.

The R packages that we used were `tidyverse` and `ggpubr` for data visualization, `lmtest` for linear regression diagnostics, `caret` for linear regression training, and `caTools` for data analysis and manipulation.

2 Data Description

We used the `Advertising.csv` dataset from Trevor Hastie’s “An Introduction to Statistical Learning” GitHub page¹. The data contains two-hundred samples of advertised products and three features (`TV`, `radio`, and `newspaper`) and one variable (`sales`) for each product. Variables `TV`, `radio`, and `newspaper` indicates the amount spent on the advertising budget for TV, radio, and newspaper, respectively, and `sales` indicates the number of sales of the product². `TV` and `radio` are slightly skew right, while `sales` is normally distributed. `newspaper` is extremely skew right.

We can see from figure 1 that the `newspaper` column clearly has an outlier in the histogram, which we remove. The data adequately passed all our tests of linearity (See Analysis Section), so no additional pruning was necessary.

¹<https://trevorhastie.github.io/ISLR/Advertising.csv>

²<https://search.r-project.org/CRAN/refmans/glmtoolbox/html/advertising.html>

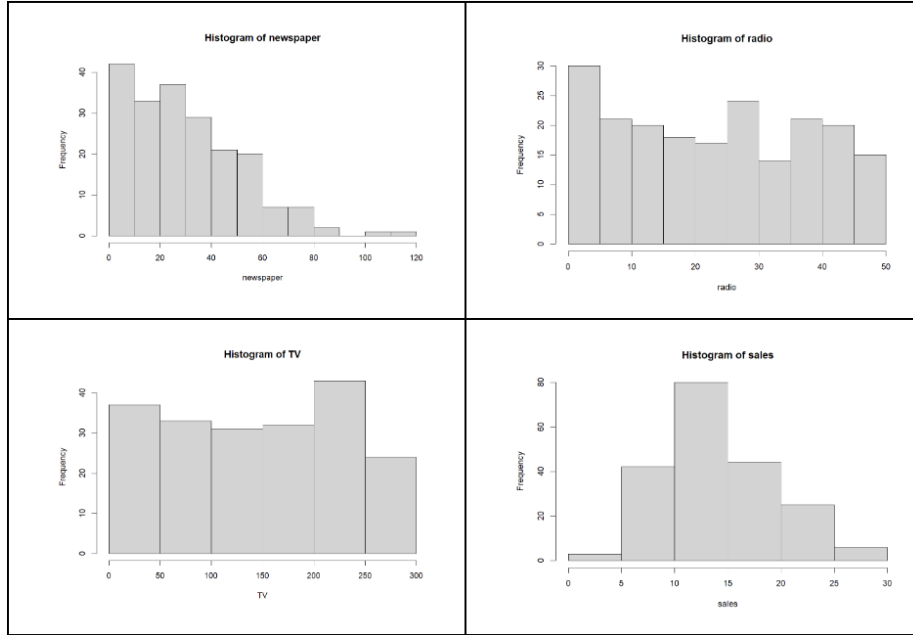


Figure 1: Histograms representing the frequency of different data values for each independent variable

3 Analysis

We used 80% of our data to train the model and the remaining 20% was used to test the model. We used the training set to train a multiple linear regression model with TV and radio as predictors, and sales as the output. We created a diagonal plot of this model (figure 2), which somewhat satisfies the linearity constraints.

The residuals have approximately mean zero, which satisfies our linearity assumption. The Durbin Watson test gives us a P value of 0.71, which is greater than 0.05, so we fail to reject the null hypothesis that the predictors are linearly independent. This is also supported by figure 3, which shows low correlation between predictors. The residuals vs fitted plot shows that the linear model doesn't fully represent the relationship between advertising budget (TV, radio) and sales. The residual errors also do not have a mean value of zero. The residual errors have approximately constant variance, since the scale-location plot does not exhibit any patterns. The Q-Q residuals show a very slight non-linear pattern, which means our assumption of normal residuals is not completely correct, but very close. The residuals vs leverage plot shows points with higher leverage also have high residuals which is not good for our model because the high influence points do not fit the model well and have high impact on the model.

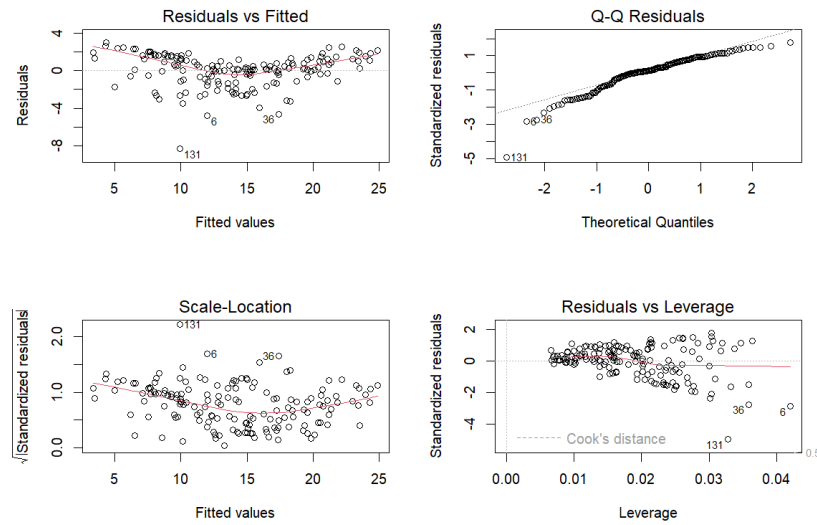


Figure 2: Diagonal plots demonstrating validity of assumptions needed for linear model

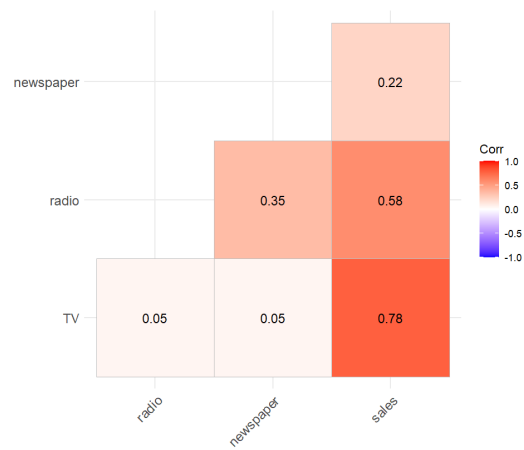


Figure 3: Correlation matrix between variables

The code used to perform analysis, as well as its output, can be found in Appendix A and B respectively.

4 Model Evaluation

We begin our model evaluation by analyzing the performance of Model 1, which includes all three predictors: TV, radio, and newspaper. Below is the summary table for Model 1:

4.1 Model 1 (Full Model)

- Predictors: **TV, radio, and newspaper**

- Equation:

$$\hat{y} = \beta_0 + \beta_1 TV + \beta_2 radio + \beta_3 newspaper \quad (1)$$

- Adjusted R^2 : **0.8915**
- MSE on test set: **2.08497**
- Coefficients and statistical significance:

Predictor	Coefficient	Std. Error	t-value	p-value
Intercept	2.5064	0.3863	6.489	1.12e-09
TV	0.0472	0.0016	29.832	<2e-16
Radio	0.1926	0.0102	18.956	<2e-16
Newspaper	-0.0017	0.0072	-0.233	0.816 (not significant)

As observed, the p-value for the newspaper predictor is 0.816, which is much greater than 0.05. This indicates that newspaper advertising does not significantly contribute to predicting sales and should be removed from the model. Consequently, we eliminate the newspaper variable and define our refined model, Model 2, which includes only TV and radio as predictors.

4.2 Model 2 (Reduced Model using Subset Selection)

- Predictors: **TV and radio**

- Equation:

$$\hat{y} = \beta_0 + \beta_1 TV + \beta_2 radio \quad (2)$$

- Adjusted R^2 : **0.8921** (higher than Model 1, despite one fewer predictor)
- MSE on test set: **2.08459**
- Coefficients and statistical significance:

Predictor	Coefficient	Std. Error	t-value	p-value
Intercept	2.4734	0.3584	6.902	1.24e-10
TV	0.0472	0.0016	29.931	<2e-16
Radio	0.1919	0.0097	19.819	<2e-16

To confirm the superiority of Model 2 over Model 1, we compare their performance metrics. The adjusted R^2 value for Model 1 is 0.8915, whereas for Model 2, it is 0.8921. Since the adjusted R^2 value slightly increases in Model 2, this indicates a better fit by removing the newspaper predictor. Furthermore, the Mean Squared Error (MSE) values for the test set are:

- Model 1: 2.08497
- Model 2: 2.08459

Since Model 2 achieves a marginally lower MSE and a higher adjusted R^2 , we conclude that it is the better model for predicting sales. Therefore, Model 2 is selected as our final model for making predictions.

4.3 Durbin-Watson Test

- The test for autocorrelation yielded a **DW statistic of 1.8874** and a **p-value of 0.2385**.
- Conclusion: We fail to reject the null hypothesis, meaning **no significant autocorrelation in residuals**, which supports model validity.

4.4 Residual Analysis

- **Histogram and Q-Q plot of residuals** suggest they are approximately normally distributed, though with slight deviation.
- **Residuals vs. Fitted plot** does not show clear heteroscedasticity, meaning variance is fairly constant.
- **Cook's Distance test for influential points** found several high-leverage points (e.g., indices 3, 6, 26, 36, 76, 77, 92, 118, 127, 131, 179), but they were not removed as their impact was minor.

5 Conclusion

Overall we saw very strong performance from our final model, with an adjusted R^2 value of 0.8962 indicating strong fitting to the testing data. Similarly, we saw a low MSE on testing data of 2.08, indicating accurate predictions when used on testing data.

However, according to our analysis, we find that our assumption of normal residuals may not completely hold, with our Q-Q residuals showing slight deviation from linear pattern. This means that our data may not be perfectly linear.

In the future, we may want to consider our model with additional data points in the testing set. As we only had around 200 data points, we were only able to split our data 80-20 testing training, and it may be more beneficial to have a more even split for additional testing data, but that would require having access to more data points.

References

- [1] Hastie, T. (n.d.). *Advertising.csv* [Dataset]. <https://trevorhastie.github.io/ISLR/Advertising.csv>
- [2] *Multiple Linear regression in R - articles - STHDA*. (2018, October 3). <http://www.sthda.com/english/articles/40-regression-analysis/168-multiple-linear-regression-in-r/>
- [3] *R: Advertising*. (n.d.). <https://search.r-project.org/CRAN/refmans/glmtoolbox/html/advertising.html>

6 Appendices

6.1 Appendix A - Source Code

```
if (!requireNamespace("car", quietly = TRUE)) {
  install.packages("car")
}

install.packages("tidyverse")
install.packages("caTools")
install.packages("ggcorrplot")

library(car)
library(lmtest)
library(tidyverse)
library(caTools)
library(ggcorrplot)

advertising <- read.csv("Advertising.csv")
str(advertising)
head(advertising, 4)
```

```

# Visualize variable distributions
for (col_name in names(advertising)) {
  if (is.numeric(advertising[[col_name]])) {
    hist(advertising[[col_name]],
          main = paste("Histogram of", col_name),
          xlab = col_name)
  }
}

# Find the outlier in newspaper
Q1 <- quantile(advertising$newspaper, 0.25)
Q3 <- quantile(advertising$newspaper, 0.75)
IQR <- Q3 - Q1

lower_bound <- Q1 - 1.5 * IQR
upper_bound <- Q3 + 1.5 * IQR

outlier_indices <- which(advertising$newspaper < lower_bound | advertising$newspaper > upper_bound)

print(paste("Outlier value:", advertising$newspapers[outlier_indices]))
print(paste("Outlier indices:", outlier_indices))
advertising <- advertising[-outlier_indices, ]

# Train test split
set.seed(123)
split <- sample.split(advertising$sales, SplitRatio = 0.8)
train_set <- subset(advertising, split == TRUE)
test_set <- subset(advertising, split == FALSE)
dim(train_set)
dim(test_set)

model1 <- lm(sales ~ TV + radio + newspaper, data = train_set)
summary(model1)
summary(model1)$coefficients

model2 <- lm(sales ~ TV + radio, data = train_set)
summary(model2)
confint(model2)

# hist(residuals(model2), main = "Histogram of Residuals", xlab = "Residuals", col = "lightblue")
qqnorm(residuals(model2))
qqline(residuals(model2), col = "red", lwd = 2)

par(mfrow = c(2, 2)) # Arrange plots in a 2x2 grid

```

```

plot(model2)
par(mfrow = c(1, 1)) # Reset plot layout

dwtest(model2)
# sigma(model2)/mean(marketing$sales)

#find MSE on test set
predictions_model1 <- predict(model1, newdata = test_set)
predictions_model2 <- predict(model2, newdata = test_set)

mse_model1 <- mean((test_set$sales - predictions_model1)^2)
print(paste("MSE for model1:", mse_model1))

mse_model2 <- mean((test_set$sales - predictions_model2)^2)
print(paste("MSE for model2:", mse_model2))

#Cook's distance to remove influential points
cooksD <- cooks.distance(model2)
plot(cooksD,type="b",pch=18,col="red")
influential <- cooksD[(cooksD > (3 * mean(cooksD, na.rm = TRUE)))]
print("Indices of influential points:")
print(names(influential))

#correlation between predictors
reduced_data <- subset(advertising, select=-X)
corr_matrix = round(cor(reduced_data),2)

ggcorrplot(corr_matrix, hc.order = FALSE, type="lower", lab=TRUE)

```

6.2 Appendix B - Code Output

```

> dim(train_set)
[1] 158  5

> dim(test_set)
[1] 40  5

> model1 <- lm(sales ~ TV + radio + newspaper, data = train_set)
> summary(model1)

Call:
lm(formula = sales ~ TV + radio + newspaper, data = train_set)

Residuals:
    Min       1Q   Median       3Q      Max

```


-8.5519 -0.9765 0.3005 1.1592 3.2111

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	2.506380	0.386262	6.489	1.12e-09 ***
TV	0.047179	0.001582	29.832	< 2e-16 ***
radio	0.192607	0.010161	18.956	< 2e-16 ***
newspaper	-0.001688	0.007246	-0.233	0.816

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.768 on 154 degrees of freedom

Multiple R-squared: 0.8935, Adjusted R-squared: 0.8915

F-statistic: 430.8 on 3 and 154 DF, p-value: < 2.2e-16

```
> summary(model1)$coefficients
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	2.506379918	0.386261532	6.488816	1.120510e-09
TV	0.047179293	0.001581523	29.831566	6.971047e-66
radio	0.192606578	0.010160618	18.956187	4.186350e-42
newspaper	-0.001688197	0.007245947	-0.232985	8.160824e-01

```
> model2 <- lm(sales ~ TV + radio, data = train_set)
```

```
> summary(model2)
```

Call:

```
lm(formula = sales ~ TV + radio, data = train_set)
```

Residuals:

Min	1Q	Median	3Q	Max
-8.5062	-0.9427	0.3082	1.1914	3.1897

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	2.473445	0.358367	6.902	1.24e-10 ***
TV	0.047167	0.001576	29.931	< 2e-16 ***
radio	0.191912	0.009683	19.819	< 2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.763 on 155 degrees of freedom

Multiple R-squared: 0.8935, Adjusted R-squared: 0.8921

F-statistic: 650.2 on 2 and 155 DF, p-value: < 2.2e-16

```
> confint(model2)
```

	2.5 %	97.5 %
--	-------	--------

```

(Intercept) 1.76553196 3.18135798
TV          0.04405435 0.05028018
radio       0.17278341 0.21103991

> dwtest(model2)

Durbin-Watson test

data: model2
DW = 1.8874, p-value = 0.2385
alternative hypothesis: true autocorrelation is greater than 0

> print(paste("MSE for model1:", mse_model1))
[1] "MSE for model1: 2.08497017705538"

> print(paste("MSE for model2:", mse_model2))
[1] "MSE for model2: 2.08459311158263"

> print("Indices of influential points:")
[1] "Indices of influential points:"

> print(names(influential))
[1] "3" "6" "26" "36" "76" "77" "92" "118" "127" "131" "179"

```